

KNN: Bias-Variance trade-off

Victor Gallego y Roi Naveiro

09/04/2019

```
library(class)
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##      cov, smooth, var
```

En este ejercicio, entrenaremos un clasificador KNN para aprender a distinguir imágenes del dígito “8” de otras del dígito “9”. Para ello, vamos a usar las proyecciones a 2D que nos daba el análisis de componentes principales.

Funciones auxiliares

- `show_digit`: Hace una gráfica del dígito en cuestión.
- `load_image_file`: Para cargar las imágenes de los dígitos
- `load_label_file`: Para cargar las etiquetas

```
show_digit = function(arr784, col = gray(12:1 / 12), ...) {
  image(matrix(as.matrix(arr784[-785]), nrow = 28)[, 28:1], col = col, ...)
}

load_image_file = function(filename) {
  ret = list()
  f = file(filename, 'rb')
  readBin(f, 'integer', n = 1, size = 4, endian = 'big')
  n = readBin(f, 'integer', n = 1, size = 4, endian = 'big')
  nrow = readBin(f, 'integer', n = 1, size = 4, endian = 'big')
  ncol = readBin(f, 'integer', n = 1, size = 4, endian = 'big')
  x = readBin(f, 'integer', n = n * nrow * ncol, size = 1, signed = FALSE)
  close(f)
  data.frame(matrix(x, ncol = nrow * ncol, byrow = TRUE))
}

load_label_file = function(filename) {
  f = file(filename, 'rb')
  readBin(f, 'integer', n = 1, size = 4, endian = 'big')
  n = readBin(f, 'integer', n = 1, size = 4, endian = 'big')
  y = readBin(f, 'integer', n = n, size = 1, signed = FALSE)
```

```

    close(f)
  }
}

```

Lectura de Datos

Cargamos el dataset MNIST.

```

df = load_image_file("src/t10k-images.idx3-ubyte")
df$y = as.factor(load_label_file("src/t10k-labels.idx1-ubyte"))

```


Esta base de datos consta de 10000 imágenes en escala de gris a 28 x 28, de los dígitos del 0 al 9 (escritos a mano).

```
dim(df)
```

```
## [1] 10000 785
```

Visualizamos algún ejemplo

```
show_digit(df[1, ])
```



únicamente las imágenes de los dígitos 8 y 9.

```
df2 = df[df$y == '8' | df$y == '9',]
```

Creación de conjuntos de train, test y validación.

Divide los datos en train y test, utilizando porcentajes 70, 30; respectivamente.

```

size_train = floor(0.7 * nrow(df2))
#size_test = floor(0.3 * nrow(df2))
#size_val = floor(0.2 * nrow(df2))
##
ind_train = sample(1:nrow(df2), size=size_train)

```

```

train = df2[ind_train,]
test = df2[-ind_train,]

#ind_val = sample(1:nrow(test_val), size=size_val)
#validation = test_val[ind_val,]
#test = test_val[-ind_val,]

```

Proyección a 2D usando PCA

Proyecta los datos de entrenamiento a dos dimensiones usando el paquete prcomp


```

proy_pca <- prcomp(train[, 1:28~2], retx = T) ## Ojo, quitar LABEL, sino son trampas

train_proy = data.frame(proy_pca$x[, 1:2])
train_proy$label = train$y
train_proy$label = factor(train_proy$label)
levels(train_proy$label) = c("n8", "n9")

p = ggplot(train_proy, aes(x = PC1, y=PC2, colour=label)) + geom_point()
p

```



Proyecta los datos de test y validación a 2D (OJO, usa las matrices de proyección generadas por el PCA del conjunto de train, de otra manera son trampas. Piensa por qué).


```

test_proy = scale(test[, 1:28~2], proy_pca$center, proy_pca$scale) %%% proy_pca$rotation
test_proy = data.frame(test_proy[, 1:2])
#test_proy = scale(test[, 1:28~2], proy_pca$center, proy_pca$scale) %%% proy_pca$rotation

```

```
#test_proy = data.frame(test_proy)

test_proy$label = test$y
test_proy$label = factor(test_proy$label)
levels(test_proy$label) = c("n8", "n9")
p = ggplot(test_proy, aes(x = PC1, y=PC2, colour=label) ) + geom_point()
p
```



Entrenamiento

Entrena un clasificador KNN usando el paquete caret. Usar validación cruzada con 5 folds y 3 repeticiones para estimar el número óptimo de vecinos. Primero definir los controles del training.

```
# Setting up train controls
repeats = 3
numbers = 5
tune1 = 100

x = trainControl(method = "repeatedcv",
                 number = numbers,
                 repeats = repeats,
                 classProbs = TRUE,
                 summaryFunction = twoClassSummary)
```

Una vez definidos, entrenar el algoritmo

```

modell1 <- train(label ~ ., data = train_proy, method = "knn",
               trControl = x,
               preProcess = c("center", "scale"),
               metric = "ROC",
               tuneLength = tune1)

# Summary of model
modell1


## k-Nearest Neighbors
##
## 1388 samples
##    2 predictors
##    2 classes: 'n8', 'n9'
##
## Pre-processing: centered (2), scaled (2)
## Resampling: Cross-Validated (5 fold, repeated 3 times)
## Summary of sample sizes: 1111, 1110, 1109, 1111, 1111, 1111, ...
## Resampling results across tuning parameters:
##
##    k    ROC      Sens      Spec
##    5  0.9734174  0.9444467  0.9470128
##    7  0.9793111  0.9463515  0.9470128
##    9  0.9819440  0.9506749  0.9489487
##   11  0.9836568  0.9502021  0.9470128
##   13  0.9844581  0.9492395  0.9479790
##   15  0.9848738  0.9506817  0.9494213
##   17  0.9851699  0.9497191  0.9484586
##   19  0.9863185  0.9482905  0.9470128
##   21  0.9871700  0.9468482  0.9494213
##   23  0.9874824  0.9463720  0.9494179
##   25  0.9878914  0.9444536  0.9489348
##   27  0.9882906  0.9458890  0.9498940
##   29  0.9884527  0.9449298  0.9503736
##   31  0.9884425  0.9430147  0.9503701
##   33  0.9885103  0.9444536  0.9518125
##   35  0.9884703  0.9449332  0.9527752
##   37  0.9885070  0.9458924  0.9542210
##   39  0.9888660  0.9454162  0.9547006
##   41  0.9888573  0.9458959  0.9537379
##   43  0.9888260  0.9444536  0.9547006
##   45  0.9890455  0.9478040  0.9551802
##   47  0.9890560  0.9449298  0.9561464
##   49  0.9890993  0.9449366  0.9561499
##   51  0.9892281  0.9449332  0.9561464
##   53  0.9892465  0.9420589  0.9556633
##   55  0.9891966  0.9434978  0.9551837
##   57  0.9893831  0.9444570  0.9546971
##   59  0.9895125  0.9444570  0.9537344
##   61  0.9893999  0.9439774  0.9527752
##   63  0.9894326  0.9425420  0.9532583
##   65  0.9893879  0.9435012  0.9542210
##   67  0.9892927  0.9435012  0.9532583
##   69  0.9893150  0.9435012  0.9537379

```

##	71	0.9892752	0.9449400	0.9532583
##	73	0.9891629	0.9444604	0.9537414
##	75	0.9891352	0.9449400	0.9542175
##	77	0.9890730	0.9458993	0.9542140
##	79	0.9891216	0.9458993	0.9547006
##	81	0.9890904	0.9449400	0.9547006
##	83	0.9890403	0.9454197	0.9537379
##	85	0.9890144	0.9449400	0.9542175
##	87	0.9890454	0.9444639	0.9537379
##	89	0.9890533	0.9449400	0.9542210
##	91	0.9891228	0.9439808	0.9547041
##	93	0.9890444	0.9439808	0.9547041
##	95	0.9890198	0.9430250	0.9537379
##	97	0.9888986	0.9435046	0.9542175
##	99	0.9888430	0.9435046	0.9542175
##	101	0.9887931	0.9435046	0.9542210
##	103	0.9887912	0.9444639	0.9542210
##	105	0.9888123	0.9439842	0.9547006
##	107	0.9886981	0.9439808	0.9556633
##	109	0.9886791	0.9435012	0.9551837
##	111	0.9886322	0.9444604	0.9566330
##	113	0.9886843	0.9439808	0.9571160
##	115	0.9886514	0.9444604	0.9571160
##	117	0.9886634	0.9449400	0.9561499
##	119	0.9885645	0.9444639	0.9561499
##	121	0.9886108	0.9430250	0.9561499
##	123	0.9885469	0.9444639	0.9561499
##	125	0.9884619	0.9439842	0.9561499
##	127	0.9884197	0.9439842	0.9556668
##	129	0.9883922	0.9439842	0.9561499
##	131	0.9883176	0.9430250	0.9566330
##	133	0.9883105	0.9430250	0.9561499
##	135	0.9882030	0.9425454	0.9561499
##	137	0.9881528	0.9430250	0.9566330
##	139	0.9880852	0.9411065	0.9566330
##	141	0.9880035	0.9406269	0.9566330
##	143	0.9879740	0.9406269	0.9571160
##	145	0.9879028	0.9415862	0.9566330
##	147	0.9878626	0.9411100	0.9566330
##	149	0.9879077	0.9411100	0.9566330
##	151	0.9878763	0.9401507	0.9571160
##	153	0.9878744	0.9396677	0.9566330
##	155	0.9878536	0.9391881	0.9561499
##	157	0.9878550	0.9387119	0.9566330
##	159	0.9878050	0.9387119	0.9566330
##	161	0.9877459	0.9396711	0.9556702
##	163	0.9876924	0.9396711	0.9566364
##	165	0.9876564	0.9401507	0.9571160
##	167	0.9876600	0.9396711	0.9566364
##	169	0.9877883	0.9401507	0.9561533
##	171	0.9877328	0.9401507	0.9561533
##	173	0.9876983	0.9406304	0.9566330
##	175	0.9876497	0.9401507	0.9571126
##	177	0.9875510	0.9406304	0.9566330

```
## 179 0.9875112 0.9401507 0.9566330
## 181 0.9874504 0.9396711 0.9566330
## 183 0.9873864 0.9391915 0.9566330
## 185 0.9873325 0.9396711 0.9571126
## 187 0.9872790 0.9401473 0.9566330
## 189 0.9873450 0.9406269 0.9566330
## 191 0.9873155 0.9406269 0.9566330
## 193 0.9873155 0.9411065 0.9566330
## 195 0.9872894 0.9401473 0.9566330
## 197 0.9872199 0.9406269 0.9566330
## 199 0.9871158 0.9411065 0.9566330
## 201 0.9870848 0.9411065 0.9566330
## 203 0.9870639 0.9396677 0.9566330
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was k = 59.
```

```
plot(model1)
```



Representar el número de vecinos frente al valor de AUC. ¿Cuál es el número óptimo de vecinos?

```
model1$bestTune
```

```
## k
## 28 59
```

AUC en el conjunto de test

Estima el valor de la AUC en el conjunto de test y pinta la curva ROC

```

preds = predict(model1, newdata = test_proy, type = "prob")
roc_obj = roc(test_proy$label, preds[,1])
auc(roc_obj)


```

```
## Area under the curve: 0.9783
```

```

roc_full_resolution <- roc(test_proy$label, preds[,1])
plot(roc_full_resolution, print.auc=TRUE)

```



Overfitting

Juega con el valor del número de vecinos para entender el comportamiento observado en el gráfico anterior.

```

k = 95
title = paste0(k, "-nearest neighbour")
x <- train_proy[, c("PC1", "PC2")]
g <- train_proy$label
px1 <- seq(min(train_proy$PC1), max(train_proy$PC1), length.out = 20)
px2 <- seq(min(train_proy$PC2), max(train_proy$PC2), length.out = 20)
xnew <- expand.grid(px1, px2)
mod15 <- knn(x, xnew, g, k=k, prob=TRUE)
prob <- attr(mod15, "prob")
prob <- ifelse(mod15=="n8", prob, 1-prob)

prob15 <- matrix(prob, length(px1), length(px2))
par(mar=rep(2,4))
contour(px1, px2, prob15, levels=0.7, labels="", xlab="", ylab="", main=
  title, axes=FALSE)
points(x, col=ifelse(g=="n8", "coral", "cornflowerblue"))

```



```
gd <- expand.grid(x=px1, y=px2)
points(gd, pch=".", cex=3.0, col=ifelse(prob15>0.5, "coral", "cornflowerblue"))
box()
```

95-nearest neighbour

